Paluri Sai Sasidhar

9494929237 | palurisaisasidhar@gmail.com | LinkedIn | GitHub | LeetCode | CodeChef

Education

Indian Institute of Information Technology & Management, Gwalior **June 2023 – Expected May 2027**
Pursuing B.Tech in Computer Science and Engineering *Gwalior, Madhya Pradesh*

Relevant Coursework

Data Structures & Algorithms, Operating Systems, Database Management (DBMS), Computer Networks, Object-Oriented Programming, Cryptography & Network Security, Software Engineering, Design and Analysis of Algorithms.

Experience

Tecdatum Infoservices Pvt. Ltd. **May 2026 – July 2026**
AI Intern *Hyderabad, Telangana*

- **Edge AI & ANPR System:** Engineered a modular Python-based software architecture and Docker-based containerized deployment pipeline for an Automatic Number Plate Recognition (ANPR) system targeting industrial edge devices (NVIDIA Jetson Nano).
- **Hardware-Accelerated Interface:** Built a custom TensorRT inference pipeline to compile and execute YOLOv8 models, programmatically managing GPU memory allocation and asynchronous execution contexts to reduce latency by 68% while maintaining 95% accuracy.
- **Escape Route Optimization:** Designed a geospatial routing engine using OSMnx and NetworkX to model complex street grids; optimized pathfinding efficiency by developing a custom pruned Dijkstra algorithm that truncated search branches using spatial coordinates.

Projects

Multimodal Vision RAG Engine | Distributed AI Search Pipeline **July 2026**
Python, PyTorch, CLIP, BM25, PyMuPDF, Qdrant, FastAPI, vLLM, Celery, Redis, Next.js

- Developed an end-to-end multimodal GenAI pipeline combining Large Vision Models (LLaVA / Qwen-VL) and CLIP embeddings for automated visual question answering and multi-page document parsing.
- Designed hybrid retrieval (Qdrant dense vectors + BM25 keyword matching) with Reciprocal Rank Fusion (RRF, k=60), reducing context retrieval latency by 42% to 38.4 ms.
- Implemented sub-5 ms LRU/Redis query-vector caching and FastAPI WebSocket streaming to maintain sub-200 ms time-to-first-token (TTFT) for client applications.
- Engineered a distributed asynchronous task queue using Celery, Redis, and PyMuPDF for concurrent multi-page document processing, reducing end-to-end ingestion time by 38%.

Sentinel Pay **September 2026**
TypeScript, Next.js, Vercel AI SDK, Vercel AI Gateway

- Converts natural-language payment requests into structured intents and evaluates amount, recipient, rolling spend, velocity, and adversarial instructions with deterministic guardrails.
- Routes results to simulated settlement, human review, or blocking; model parsing is separate from payment execution, with deterministic fallback parsing and decision traces.

Achievements / Certifications

- Active competitive programmer with peak ratings of 1801 on CodeChef and 1787 on LeetCode.
- Achieved an All India Rank (AIR) of 8741 in JEE Mains 2023 (Top 0.8% out of 1.2M+ candidates).
- Google Cloud Gen AI Academy APAC Edition Certificate, Google Cybersecurity Professional Certificate, and IBM Cybersecurity Fundamentals.

Technical Skills

AI/ML & GenAI: PyTorch, TensorFlow, HuggingFace, LLMs, LVMs, TensorRT, vLLM, RAG Pipelines, Model Quantization
Languages: Python, C/C++, JavaScript (ES6+), SQL, Java
ML Infra & Systems: CUDA (Basics), Docker, Triton Inference Server, gRPC, Multiprocessing, Linux CLI
Frameworks & Databases: FastAPI, Flask, Next.js, PostgreSQL, Redis, SQLite, RESTful API Design
Developer Tools: Git/GitHub, TensorBoard, Model Evaluation & Benchmarking Tools

Leadership / Extracurricular

Student Activity Council (Cultural) & Sports

- **Event Management:** Coordinated the IPL Auction and Pictionary events for Aurora (annual cultural fest) and spearheaded 10+ Orientation Week events, improving student satisfaction scores by 15%.
- **Athletics:** Black Belt in Karate with 9+ years of training and 30+ tournament appearances (Gold/Silver/Bronze medalist); trained competitive swimmer.